

tsla-stock

September 13, 2024

1 By Prisca

```
[79]: # TSLA stock prediction using LSTM model
```

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from tensorflow import keras
```

```
[76]: #load the data
```

```
[2]: df=pd.read_csv(r'C:\Users\USER\Documents\dataset\TSLA.csv')
```

```
[3]: df.head()
```

```
[3]:
```

	Unnamed: 0	Date	Open	High	Low	Close	Volume
0	0	2010-06-29	1.266667	1.666667	1.169333	1.592667	281494500
1	1	2010-06-30	1.719333	2.028000	1.553333	1.588667	257806500
2	2	2010-07-01	1.666667	1.728000	1.351333	1.464000	123282000
3	3	2010-07-02	1.533333	1.540000	1.247333	1.280000	77097000
4	4	2010-07-06	1.333333	1.333333	1.055333	1.074000	103003500

```
[4]: df.shape
```

```
[4]: (3534, 7)
```

```
[5]: df.drop(['Unnamed: 0'],axis=1,inplace=True)
```

```
[6]: df.set_index(['Date'],inplace=True)
```

```
[7]: df.head() # predicting stock for 2024-7-16, stock start from 2010 to 2024 july
```

```
[7]:
```

	Open	High	Low	Close	Volume
Date					
2010-06-29	1.266667	1.666667	1.169333	1.592667	281494500
2010-06-30	1.719333	2.028000	1.553333	1.588667	257806500
2010-07-01	1.666667	1.728000	1.351333	1.464000	123282000

```
2010-07-02  1.533333  1.540000  1.247333  1.280000  77097000
2010-07-06  1.333333  1.333333  1.055333  1.074000  103003500
```

```
[8]: df.isnull().sum()
```

```
[8]: Open      0
     High      0
     Low       0
     Close     0
     Volume    0
     dtype: int64
```

```
[9]: df.duplicated().sum()
```

```
[9]: 0
```


```
[77]: # selection of target column
```

```
[78]: new_df=df['Close'].values
```

```
[11]: new_df=pd.DataFrame(new_df)
```

```
[12]: plt.figure(figsize=(15,10))
     plt.plot(new_df)
```

```
[12]: [<matplotlib.lines.Line2D at 0x1324fafdaf0>]
```



```

[80]: # data scaling

[13]: from sklearn.preprocessing import MinMaxScaler

[14]: scaler=MinMaxScaler(feature_range=(0,1))

[15]: df_scaled=scaler.fit_transform(new_df)

[16]: # split the data

[98]: split_index=int(len(df_scaled)*0.8)

[99]: train_df=df_scaled[:split_index]
      test_df=df_scaled[split_index:]

[100]: train_df.shape,test_df.shape

[100]: ((2827, 1), (707, 1))

[101]: # created sequence for both train and test data

[102]: x_train=[]
      y_train=[]

      for i in range(60,len(train_df)):
          x_train.append(train_df[i-60:i, 0])
          y_train.append(train_df[i,0])

      x_train=np.array(x_train)
      y_train=np.array(y_train)

[103]: x_test=[]
      y_test=[]

      for i in range(60,len(test_df)):
          x_test.append(test_df[i-60:i,0])
          y_test.append(test_df[i,0])

      x_test=np.array(x_test)
      y_test=np.array(y_test)

[104]: x_train.shape,x_test.shape

[104]: ((2767, 60), (647, 60))

```

```
[105]: y_train.shape,y_test.shape
```

```
[105]: ((2767,), (647,))
```

2 reshape

```
[106]: x_train=x_train.reshape(x_train.shape[0],x_train.shape[1],1)
x_test=x_test.reshape(x_test.shape[0],x_test.shape[1],1)
```

```
[107]: x_train.shape,x_test.shape
```

```
[107]: ((2767, 60, 1), (647, 60, 1))
```

3 build model

```
[29]: from keras .models import Sequential
from keras .layers import Dense,Dropout,LSTM
```

```
[30]: model=Sequential()
```

```
[31]: model.add(LSTM(units=50,return_sequences=True,input_shape=( x_train.
↪shape[1],x_train.shape[2])))
model.add(Dropout(0.2))
```

```
[32]: model.add(LSTM(units=50,return_sequences=True))
model.add(Dropout(0.2))
```

```
[33]: model.add(LSTM(units=50,return_sequences=True))
model.add(Dropout(0.2))
```

```
[34]: model.add(LSTM(units=50,return_sequences=True))
model.add(Dropout(0.2))
```

```
[35]: model.add(LSTM(units=50,return_sequences=True))
model.add(Dropout(0.2))
```

```
[36]: model.add(LSTM(units=50,return_sequences=False))
model.add(Dropout(0.2))
```

```
[37]: model.add(Dense(units=1))
```

```
[ ]: # compile
```

```
[39]: model.compile(optimizer='adam',loss='mean_squared_error')
```

```
[40]: #fit the model
```

```
[41]: history=model.  
      ↪fit(x_train,y_train,validation_data=(x_test,y_test),epochs=20,batch_size=2)
```

```
Epoch 1/20  
1384/1384 [=====] - 300s 202ms/step - loss: 0.0024 -  
val_loss: 0.0149  
Epoch 2/20  
1384/1384 [=====] - 281s 203ms/step - loss: 0.0014 -  
val_loss: 0.0110  
Epoch 3/20  
1384/1384 [=====] - 279s 201ms/step - loss: 0.0012 -  
val_loss: 0.0102  
Epoch 4/20  
1384/1384 [=====] - 273s 197ms/step - loss: 0.0010 -  
val_loss: 0.0100  
Epoch 5/20  
1384/1384 [=====] - 278s 201ms/step - loss: 9.9633e-04  
- val_loss: 0.0259  
Epoch 6/20  
1384/1384 [=====] - 283s 205ms/step - loss: 8.8110e-04  
- val_loss: 0.0090  
Epoch 7/20  
1384/1384 [=====] - 274s 198ms/step - loss: 8.3555e-04  
- val_loss: 0.0090  
Epoch 8/20  
1384/1384 [=====] - 281s 203ms/step - loss: 7.4407e-04  
- val_loss: 0.0094  
Epoch 9/20  
1384/1384 [=====] - 274s 198ms/step - loss: 8.6287e-04  
- val_loss: 0.0077  
Epoch 10/20  
1384/1384 [=====] - 272s 196ms/step - loss: 6.2938e-04  
- val_loss: 0.0051  
Epoch 11/20  
1384/1384 [=====] - 359s 260ms/step - loss: 6.6233e-04  
- val_loss: 0.0186  
Epoch 12/20  
1384/1384 [=====] - 524s 379ms/step - loss: 7.7214e-04  
- val_loss: 0.0138  
Epoch 13/20  
1384/1384 [=====] - 809s 584ms/step - loss: 6.7801e-04  
- val_loss: 0.0058  
Epoch 14/20  
1384/1384 [=====] - 957s 691ms/step - loss: 6.4346e-04  
- val_loss: 0.0069
```

```

Epoch 15/20
1384/1384 [=====] - 897s 648ms/step - loss: 7.0937e-04
- val_loss: 0.0085
Epoch 16/20
1384/1384 [=====] - 843s 609ms/step - loss: 6.1364e-04
- val_loss: 0.0071
Epoch 17/20
1384/1384 [=====] - 929s 671ms/step - loss: 5.9391e-04
- val_loss: 0.0055
Epoch 18/20
1384/1384 [=====] - 942s 681ms/step - loss: 6.0241e-04
- val_loss: 0.0068
Epoch 19/20
1384/1384 [=====] - 383s 276ms/step - loss: 5.7224e-04
- val_loss: 0.0087
Epoch 20/20
1384/1384 [=====] - 281s 203ms/step - loss: 5.4193e-04
- val_loss: 0.0087

```


```
[90]: #model evaluation
```

```
[108]: loss_t=(history.history['loss'])
loss_v=(history.history['val_loss'])
```

```
[109]: plt.plot(loss_t)
plt.plot(loss_v)
plt.xlabel('loss')
plt.ylabel('val_loss')
plt.legend() # as the val loss is going down, the loss is low too
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

```
[109]: <matplotlib.legend.Legend at 0x1325c6f8a00>
```



4 prediction

```
[110]: predictions=model.predict(x_test)
```

```
21/21 [=====] - 2s 88ms/step
```

```
[111]: predictions=scaler.inverse_transform(predictions)
```

```
[112]: predictions.shape
```

```
[112]: (647, 1)
```

```
[113]: # Reshape y_test
```

```
[114]: y_test=y_test.reshape(-1,1)
```

```
[115]: y_test.shape
```

```
[115]: (647, 1)
```

```
[116]: y_test_sc=scaler.inverse_transform(y_test)
```


```
[117]: y_test_sc.shape,predictions.shape
```

```
[117]: ((647, 1), (647, 1))
```

5 stock visualization

```
[119]: plt.plot(y_test_sc)
plt.plot(predictions,color='red') # trend highly copied
```

```
[119]: [<matplotlib.lines.Line2D at 0x13253b46310>]
```



```
[56]: #evaluation
```

```
[120]: from sklearn.metrics import mean_squared_error
```

```
[121]: rmse=np.sqrt(mean_squared_error(y_test_sc,predictions))
```

```
[122]: rmse # the error is high but what matter is the trend is quite captured,if i
      ↪ train it more, the error will reduce
```



```
[122]: 38.148803216812006
```

6 lets predict the next day that is 2024-7-16

```
[123]: # reload the data
```

```
[124]: df=pd.read_csv(r'C:\Users\USER\Documents\dataset\TSLA.csv')
```

```
[125]: df.tail()
```

```
[125]:
```

	Unnamed: 0	Date	Open	High	Low	Close	\
3529	3529	2024-07-09	251.000000	265.609985	250.300003	262.329987	
3530	3530	2024-07-10	262.799988	267.589996	257.859985	263.260010	
3531	3531	2024-07-11	263.299988	271.000000	239.649994	241.029999	
3532	3532	2024-07-12	235.800003	251.839996	233.089996	248.229996	
3533	3533	2024-07-15	255.964996	265.579987	251.729996	252.639999	

	Volume
3529	160210900
3530	128519400
3531	221707300
3532	155694400
3533	142831728

```
[87]: # chose the target variable
```

```
[126]: new_df=df['Close']
```

```
[127]: # select the last 60 days
```

```
[128]: last_60_days=new_df[-60:]
```

```
[129]: last=pd.DataFrame(last_60_days)
```

```
[130]: last.shape
```

```
[130]: (60, 1)
```

```
[131]: last_days=scaler.transform(last)
```

```
C:\Users\USER\anaconda3\envs\tf_env\lib\site-packages\sklearn\base.py:458:  
UserWarning: X has feature names, but MinMaxScaler was fitted without feature  
names  
warnings.warn(  

```

```
[89]: #created sequence and append the 60 days
```

```
[132]: x_test=[]  
x_test.append(last_days)  
  
x_test=np.array(x_test)
```

```
[133]: x_test.shape
```

```
[133]: (1, 60, 1)
```

```
[134]: #reshape
```

```
[135]: x_predict=model.predict(x_test)
```

```
1/1 [=====] - 0s 71ms/step
```

```
[136]: x_predict=scaler.inverse_transform(x_predict)
```

```
[138]: x_predict #stock price for 2024-7-16
```

```
[138]: array([[223.94717]], dtype=float32)
```

```
[ ]:
```

```
[ ]:
```